

## Overview

Unsupervised learning discovers structure in unlabeled data. Instead of predicting known targets, models infer clusters, latent dimensions, or anomaly patterns.

## Core Definitions

- **Clustering:** group samples with high intra-group similarity and low inter-group similarity.
- **Dimensionality reduction:** map high-dimensional data to lower-dimensional representation while preserving useful structure.

## Clustering Methods

### k-Means

Optimizes within-cluster sum of squares:

$$\min_{\{C_k\}} \sum_{k=1}^K \sum_{x_i \in C_k} \|x_i - \mu_k\|^2$$

Fast and popular, but assumes roughly spherical clusters.

### Hierarchical Clustering

Builds cluster tree (dendrogram) via agglomerative merges or divisive splits.

### DBSCAN

Density-based clustering with parameters **eps** and **min\_samples**. Finds arbitrary-shaped clusters and labels noise.

## Dimensionality Reduction

### PCA

Linear projection maximizing explained variance.

### t-SNE and UMAP (High-Level)

Non-linear methods for visualization and neighborhood preservation; useful exploratory tools, not always stable for downstream numeric decisions.

## Distance and Similarity

Choice of metric (Euclidean, cosine, Manhattan) changes clustering behavior. Standardization usually required before distance-based methods.

## Common Pitfalls

- Forcing k-means on non-spherical data.
- Choosing **k** only by visual preference.
- Misusing t-SNE map distances as true global geometry.

# Business Case Studies & Exceptions

## Case 1: Customer Segmentation

Retail team clustered purchase behavior to tailor campaigns. Segment quality degraded when seasonal behavior changed.

Mitigation: - Refit clusters periodically. - Track silhouette and segment drift metrics.

## Case 2: Anomaly Detection via Clustering Distance

Using distance-to-centroid flagged anomalies, but scaling changes caused false alarms.

Mitigation: - Feature scaling pipeline lock. - Segment-aware thresholds and human review loop.

# Interview Questions & Answers

1. **Q: Difference between clustering and classification?**  
**A:** Clustering discovers groups without labels; classification predicts predefined labels.
2. **Q: How does k-means work?**  
**A:** Alternate assignment to nearest centroid and centroid recomputation until convergence.
3. **Q: How choose k?**  
**A:** Use elbow/silhouette/domain constraints and stability checks.
4. **Q: When use DBSCAN?**  
**A:** When clusters are arbitrary-shaped and noise detection matters.
5. **Q: Why standardize features before clustering?**  
**A:** Prevent high-scale features from dominating distance calculations.
6. **Q: PCA purpose in ML workflow?**  
**A:** Compress features, remove redundancy, and aid visualization.
7. **Q: Limitation of t-SNE?**  
**A:** Great for visualization but poor for preserving global distances/cluster sizes.
8. **Q: Silhouette score interpretation?**  
**A:** Higher means better separation/cohesion, but not sole truth metric.
9. **Q: Unsupervised output validation challenge?**  
**A:** No labels, so rely on internal metrics and business usefulness tests.
10. **Q: Business risk in unsupervised segmentation?**  
**A:** Spurious clusters can drive bad targeting or pricing decisions.